

# LoanTap Logistic Regression

## Context:

LoanTap is an online platform committed to delivering customized loan products to millennials. They innovate in an otherwise dull loan segment, to deliver instant, flexible loans on consumer friendly terms to salaried professionals and businessmen.

The data science team at LoanTap is building an underwriting layer to determine the creditworthiness of MSMEs as well as individuals.

LoanTap deploys formal credit to salaried individuals and businesses 4 main financial instruments:

Personal Loan EMI Free Loan Personal Overdraft Advance Salary Loan This case study will focus on the underwriting process behind Personal Loan only

## Problem Statement:

Given a set of attributes for an Individual, determine if a credit line should be extended to them. If so, what should the repayment terms be in business recommendations?

In [89]:

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 # Linear and Logistic regression Library
7 from sklearn.linear_model import LinearRegression, Ridge, Lasso, LogisticRegression
8 from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
9 from sklearn.model_selection import train_test_split, GridSearchCV, KFold
10 from sklearn.preprocessing import StandardScaler, MinMaxScaler, PolynomialFeatures
11 from sklearn.pipeline import make_pipeline
12 from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, precision_score, recall_score, f1_score,roc_curve, roc_auc_score
13 from sklearn.metrics import precision_recall_curve, auc
14 #stats model Library
15 import statsmodels.api as sm
16 from statsmodels.stats.outliers_influence import variance_inflation_factor
17
18 # hypothesis testing Library
19 from scipy.stats import shapiro
20
21 # math Library
22 import math
23
24 # target encoder
25 from category_encoders import TargetEncoder
26 from sklearn.preprocessing import LabelEncoder
27
28 # date
29 from datetime import date
30
31 #PCA
32 from sklearn.decomposition import PCA
33
34 #SMOTE - Balancing data
35 from imblearn.over_sampling import SMOTE
```

In [50]:

```
1 # Load Jamboree education data
2 original = pd.read_csv(r'D:\PY\course\course material\Module 13 - Intro to ML and NN\Business Case LoanTap Logistic Regression\logistic_regression.csv')
```

In [51]:

```
1 df = original.copy()
```

## Target / dependent feature - 'loan\_status'

In [52]:

```
1 df['loan_status'].unique()
```

Out[52]: array(['Fully Paid', 'Charged Off'], dtype=object)

In [53]:

```
1 df
```

Out[53]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ownership	annual_inc	...	open_acc	pub_rec	revol_bal	revol_util	total_acc	initial_list_status	application_type	mort_acc
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	RENT	117000.0	...	16.0	0.0	36369.0	41.8	25.0	w	INDIVIDUAL	0.0
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	MORTGAGE	65000.0	...	17.0	0.0	20131.0	53.3	27.0	f	INDIVIDUAL	3.0
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	RENT	43057.0	...	13.0	0.0	11987.0	92.2	26.0	f	INDIVIDUAL	0.0
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years	RENT	54000.0	...	6.0	0.0	5472.0	21.5	13.0	f	INDIVIDUAL	0.0
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years	MORTGAGE	55000.0	...	13.0	0.0	24584.0	69.8	43.0	f	INDIVIDUAL	1.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
396025	10000.0	60 months	10.99	217.38	B	B4	licensed bankere	2 years	RENT	40000.0	...	6.0	0.0	1990.0	34.3	23.0	w	INDIVIDUAL	0.0
396026	21000.0	36 months	12.29	700.42	C	C1	Agent	5 years	MORTGAGE	110000.0	...	6.0	0.0	43263.0	95.7	8.0	f	INDIVIDUAL	1.0
396027	5000.0	36 months	9.99	161.32	B	B1	City Carrier	10+ years	RENT	56500.0	...	15.0	0.0	32704.0	66.9	23.0	f	INDIVIDUAL	0.0
396028	21000.0	60 months	15.31	503.02	C	C2	Gracon Services, Inc	10+ years	MORTGAGE	64000.0	...	9.0	0.0	15704.0	53.8	20.0	f	INDIVIDUAL	5.0
396029	2000.0	36 months	13.61	67.98	C	C2	Internal Revenue Service	10+ years	RENT	42996.0	...	3.0	0.0	4292.0	91.3	19.0	f	INDIVIDUAL	NaN

396030 rows × 27 columns

In [54]:

```
1 df.columns
```

Out[54]: Index(['loan\_amnt', 'term', 'int\_rate', 'installment', 'grade', 'sub\_grade', 'emp\_title', 'emp\_length', 'home\_ownership', 'annual\_inc', 'verification\_status', 'issue\_d', 'loan\_status', 'purpose', 'title', 'dti', 'earliest\_cr\_line', 'open\_acc', 'pub\_rec', 'revol\_bal', 'revol\_util', 'total\_acc', 'initial\_list\_status', 'application\_type', 'mort\_acc', 'pub\_rec\_bankruptcies', 'address'], dtype='object')

```
In [55]: 1 df[:2].T

Out[55]:
```

	0	1
loan_amnt	10000.0	8000.0
term	36 months	36 months
int_rate	11.44	11.99
installment	329.48	265.68
grade	B	B
sub_grade	B4	B5
emp_title	Marketing	Credit analyst
emp_length	10+ years	4 years
home_ownership	RENT	MORTGAGE
annual_inc	117000.0	65000.0
verification_status	Not Verified	Not Verified
issue_d	Jan-2015	Jan-2015
loan_status	Fully Paid	Fully Paid
purpose	vacation	debt_consolidation
title	Vacation	Debt consolidation
dti	26.24	22.05
earliest_cr_line	Jun-1990	Jul-2004
open_acc	16.0	17.0
pub_rec	0.0	0.0
revol_bal	36369.0	20131.0
revol_util	41.8	53.3
total_acc	25.0	27.0
initial_list_status	w	f
application_type	INDIVIDUAL	INDIVIDUAL
mort_acc	0.0	3.0
pub_rec_bankruptcies	0.0	0.0
address	0174 Michelle Gateway\r\nMendozaberg, OK 22690 1076 Carney Fort Apt. 347\r\nLoganmouth, SD 05113	

Null values are present in below mentioned columns

'emp\_title', 'emp\_length', 'title', 'revol\_util', 'mort\_acc', 'pub\_rec\_bankruptcies'

```
In [56]: 1 df.isna().sum()

Out[56]: loan_amnt      0
term      0
int_rate  0
installment  0
grade     0
sub_grade 0
emp_title 22927
emp_length 18301
home_ownership  0
annual_inc  0
verification_status  0
issue_d      0
loan_status  0
purpose     0
title      1755
dti         0
earliest_cr_line  0
open_acc    0
pub_rec     0
revol_bal   0
revol_util  276
total_acc   0
initial_list_status  0
application_type  0
mort_acc    37795
pub_rec_bankruptcies  535
address     0
dtype: int64
```

Data are present in different scale, need to normalize the data before Logistic Regression

```
In [57]: 1 df.describe()

Out[57]:
```

	loan_amnt	int_rate	installment	annual_inc	dti	open_acc	pub_rec	revol_bal	revol_util	total_acc	mort_acc	pub_rec_bankruptcies
count	396030.000000	396030.000000	396030.000000	3.960300e+05	396030.000000	396030.000000	396030.000000	3.960300e+05	395754.000000	396030.000000	358235.000000	395495.000000
mean	14113.888089	13.639400	431.849698	7.420318e+04	17.379514	11.311153	0.178191	1.584454e+04	53.791749	25.414744	1.813991	0.121648
std	8357.441341	4.472157	250.727790	6.163762e+04	18.019092	5.137649	0.530671	2.059184e+04	24.452193	11.886991	2.147930	0.356174
min	500.000000	5.320000	16.080000	0.000000e+00	0.000000	0.000000	0.000000	0.000000e+00	0.000000	2.000000	0.000000	0.000000
25%	8000.000000	10.490000	250.330000	4.500000e+04	11.280000	8.000000	0.000000	6.025000e+03	35.800000	17.000000	0.000000	0.000000
50%	12000.000000	13.330000	375.430000	6.400000e+04	16.910000	10.000000	0.000000	1.118100e+04	54.800000	24.000000	1.000000	0.000000
75%	20000.000000	16.490000	567.300000	9.000000e+04	22.980000	14.000000	0.000000	1.962000e+04	72.900000	32.000000	3.000000	0.000000
max	40000.000000	30.990000	1533.810000	8.706582e+06	9999.000000	90.000000	86.000000	1.743266e+06	892.300000	151.000000	34.000000	8.000000

There are 27 columns, Where most datatype are Float64 and Object.

- Data type change is required
- 396030 records are availalbe in data

```
In [58]: 1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt              396030 non-null  float64
1   term                  396030 non-null  object
2   int_rate              396030 non-null  float64
3   installment           396030 non-null  float64
4   grade                 396030 non-null  object
5   sub_grade             396030 non-null  object
6   emp_title              373103 non-null  object
7   emp_length            377729 non-null  object
8   home_ownership        396030 non-null  object
9   annual_inc            396030 non-null  float64
10  verification_status   396030 non-null  object
11  issue_d               396030 non-null  object
12  loan_status           396030 non-null  object
13  purpose               396030 non-null  object
14  title                 394275 non-null  object
15  dti                   396030 non-null  float64
16  earliest_cr_line      396030 non-null  object
17  open_acc              396030 non-null  float64
18  pub_rec               396030 non-null  float64
19  revol_bal             396030 non-null  float64
20  revol_util            395754 non-null  float64
21  total_acc             396030 non-null  float64
22  initial_list_status   396030 non-null  object
23  application_type      396030 non-null  object
24  mort_acc              358235 non-null  float64
25  pub_rec_bankruptcies  395495 non-null  float64
26  address               396030 non-null  object
dtypes: float64(12), object(15)
memory usage: 81.6+ MB
```

Missing value treatment for below columns - replace by mean value

- 'emp\_length','revol\_util','mort\_acc','pub\_rec\_bankruptcies'

```
In [59]: 1 df1 = df.copy()

In [60]: 1 df1[['emp_length', 'revol_util', 'mort_acc', 'pub_rec_bankruptcies']].isna().sum()

Out[60]: emp_length      18301
revol_util         276
mort_acc           37795
pub_rec_bankruptcies  535
dtype: int64

In [61]: 1 column = ['revol_util', 'mort_acc', 'pub_rec_bankruptcies']
2 for i in column:
3     mean = df1[i].mean()
4     df1[i].fillna(value = mean, inplace=True)

In [62]: 1 df1['emp_length'].fillna(value = df1['emp_length'].mode()[0], inplace=True)
2 df1['title'].fillna(value = 'unknown', inplace=True)
3 df1['emp_title'].fillna(value = 'unknown', inplace=True)
```

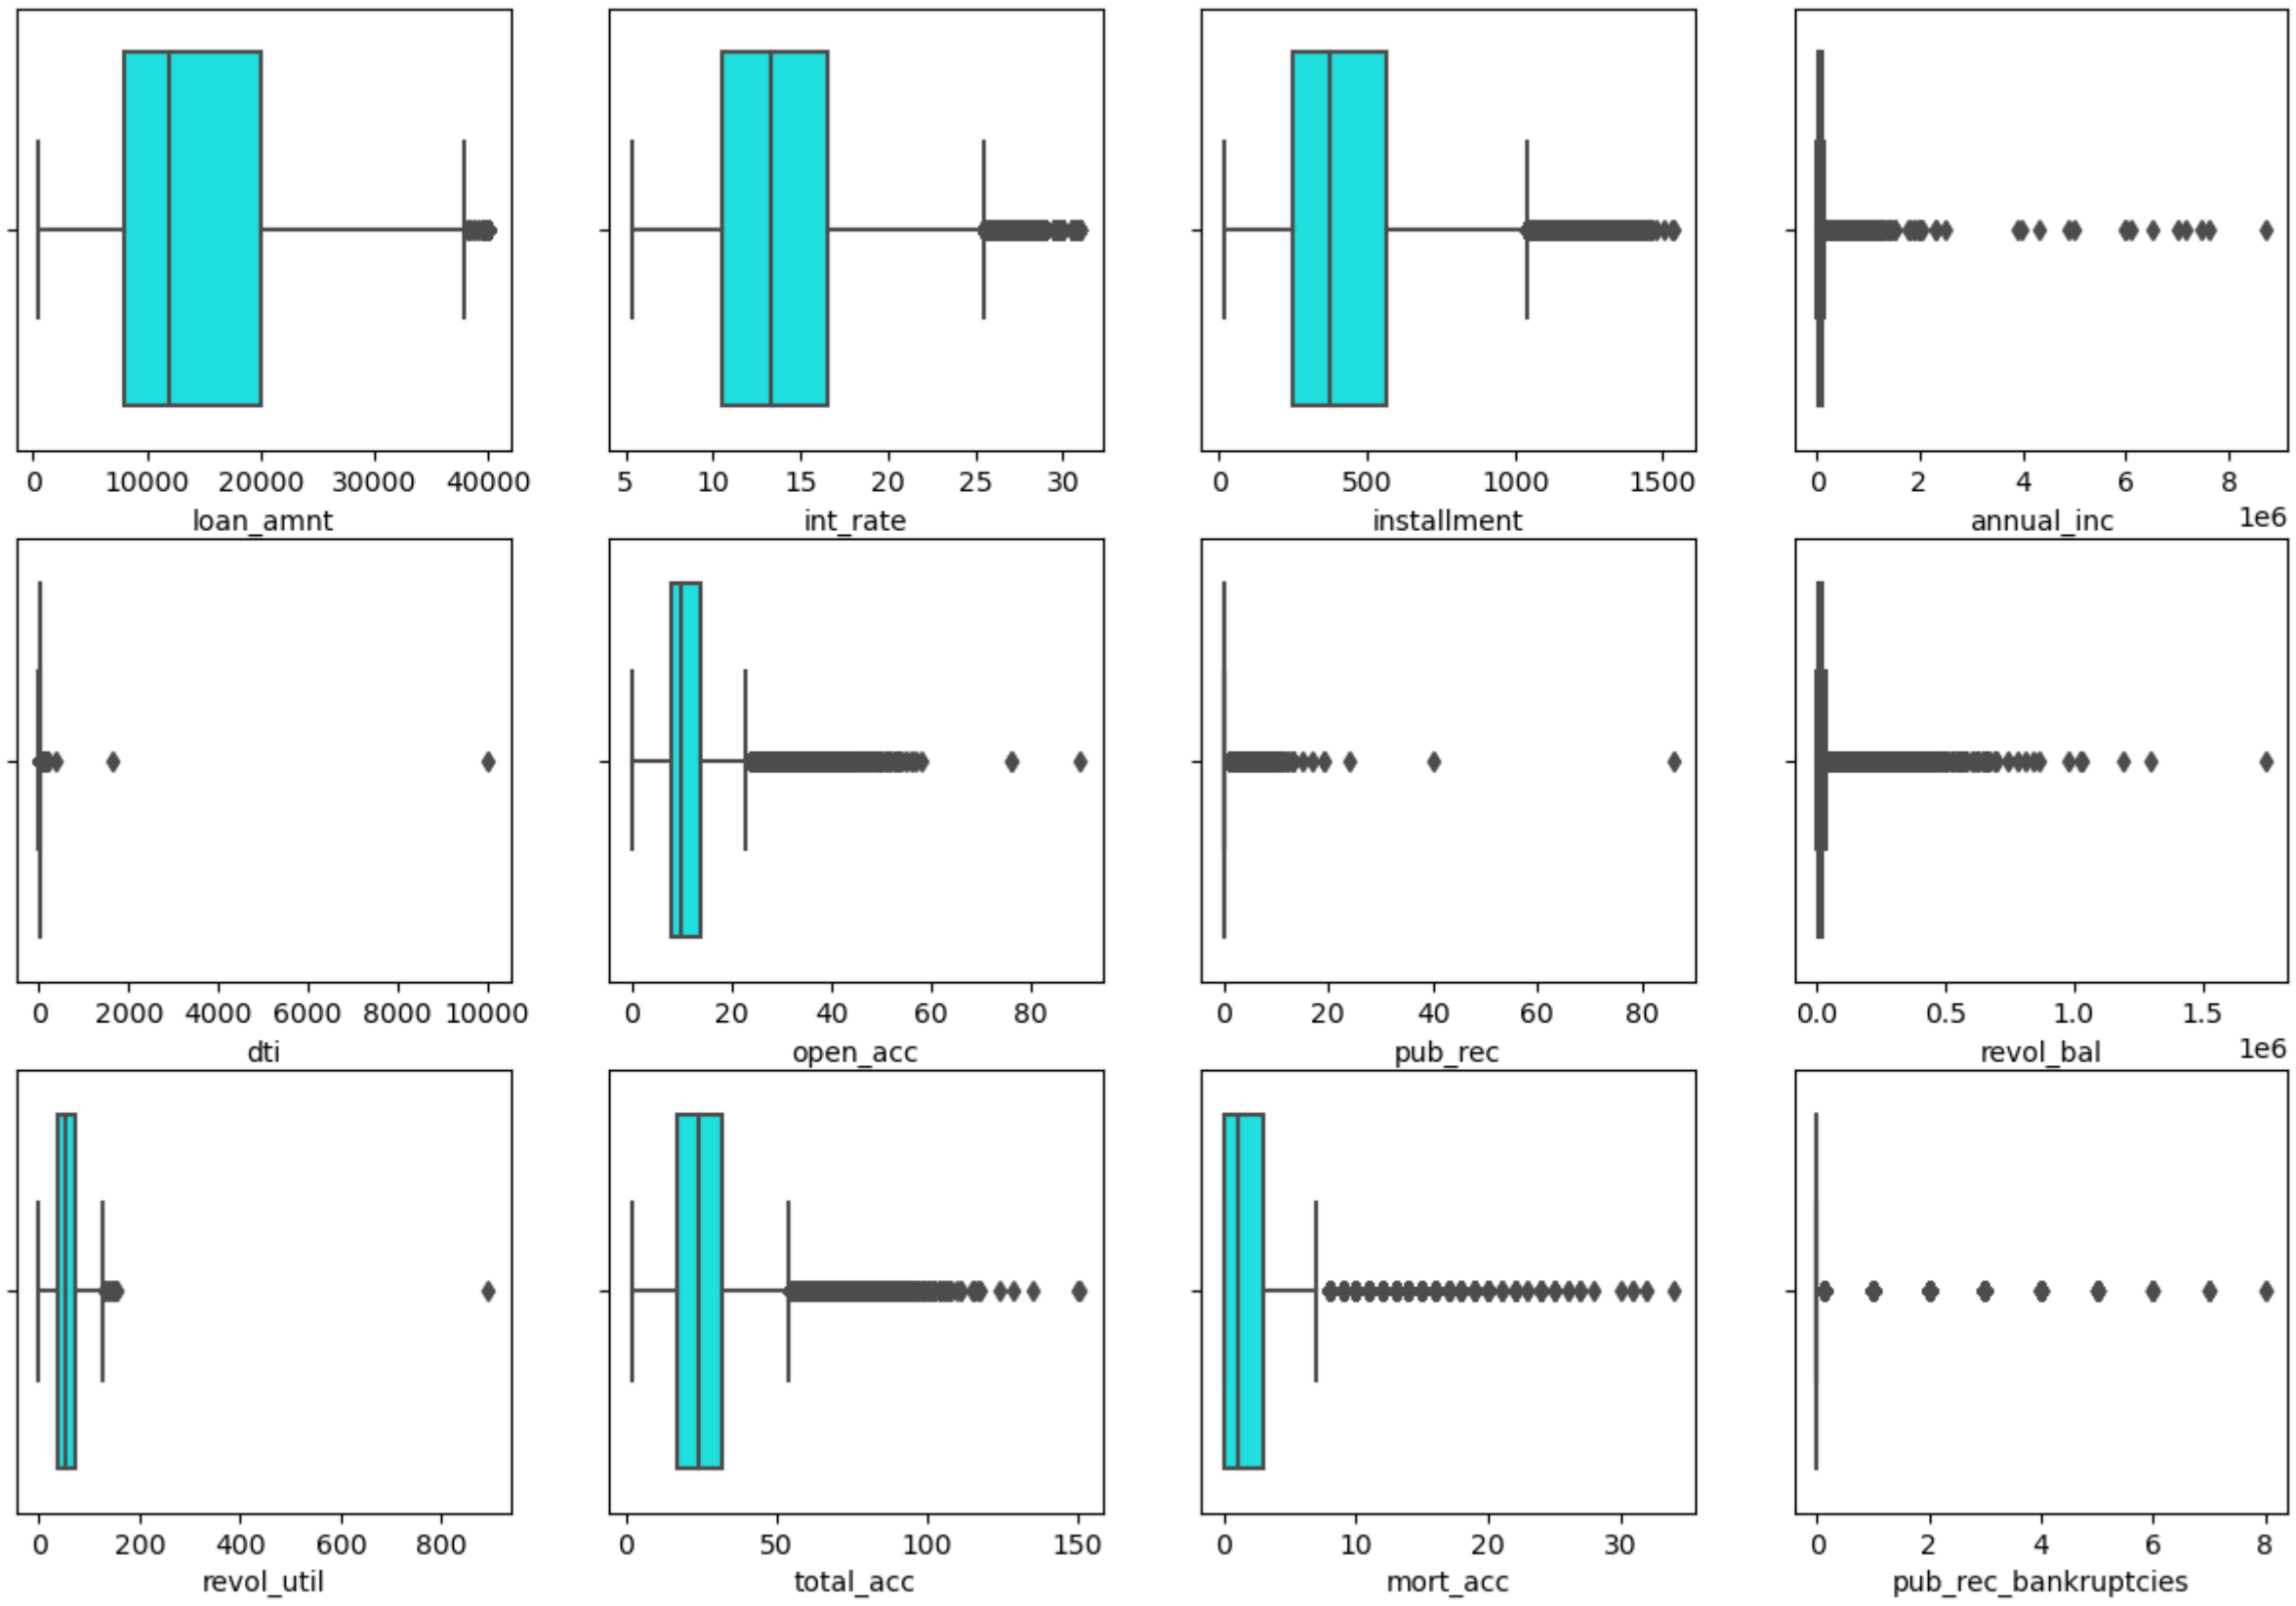
Outlier treatment

- Data as outlier, outlier to be removed by converting all outlier data with IQR range.

```
In [63]: 1 df2 = df1.copy()
```

In [64]:

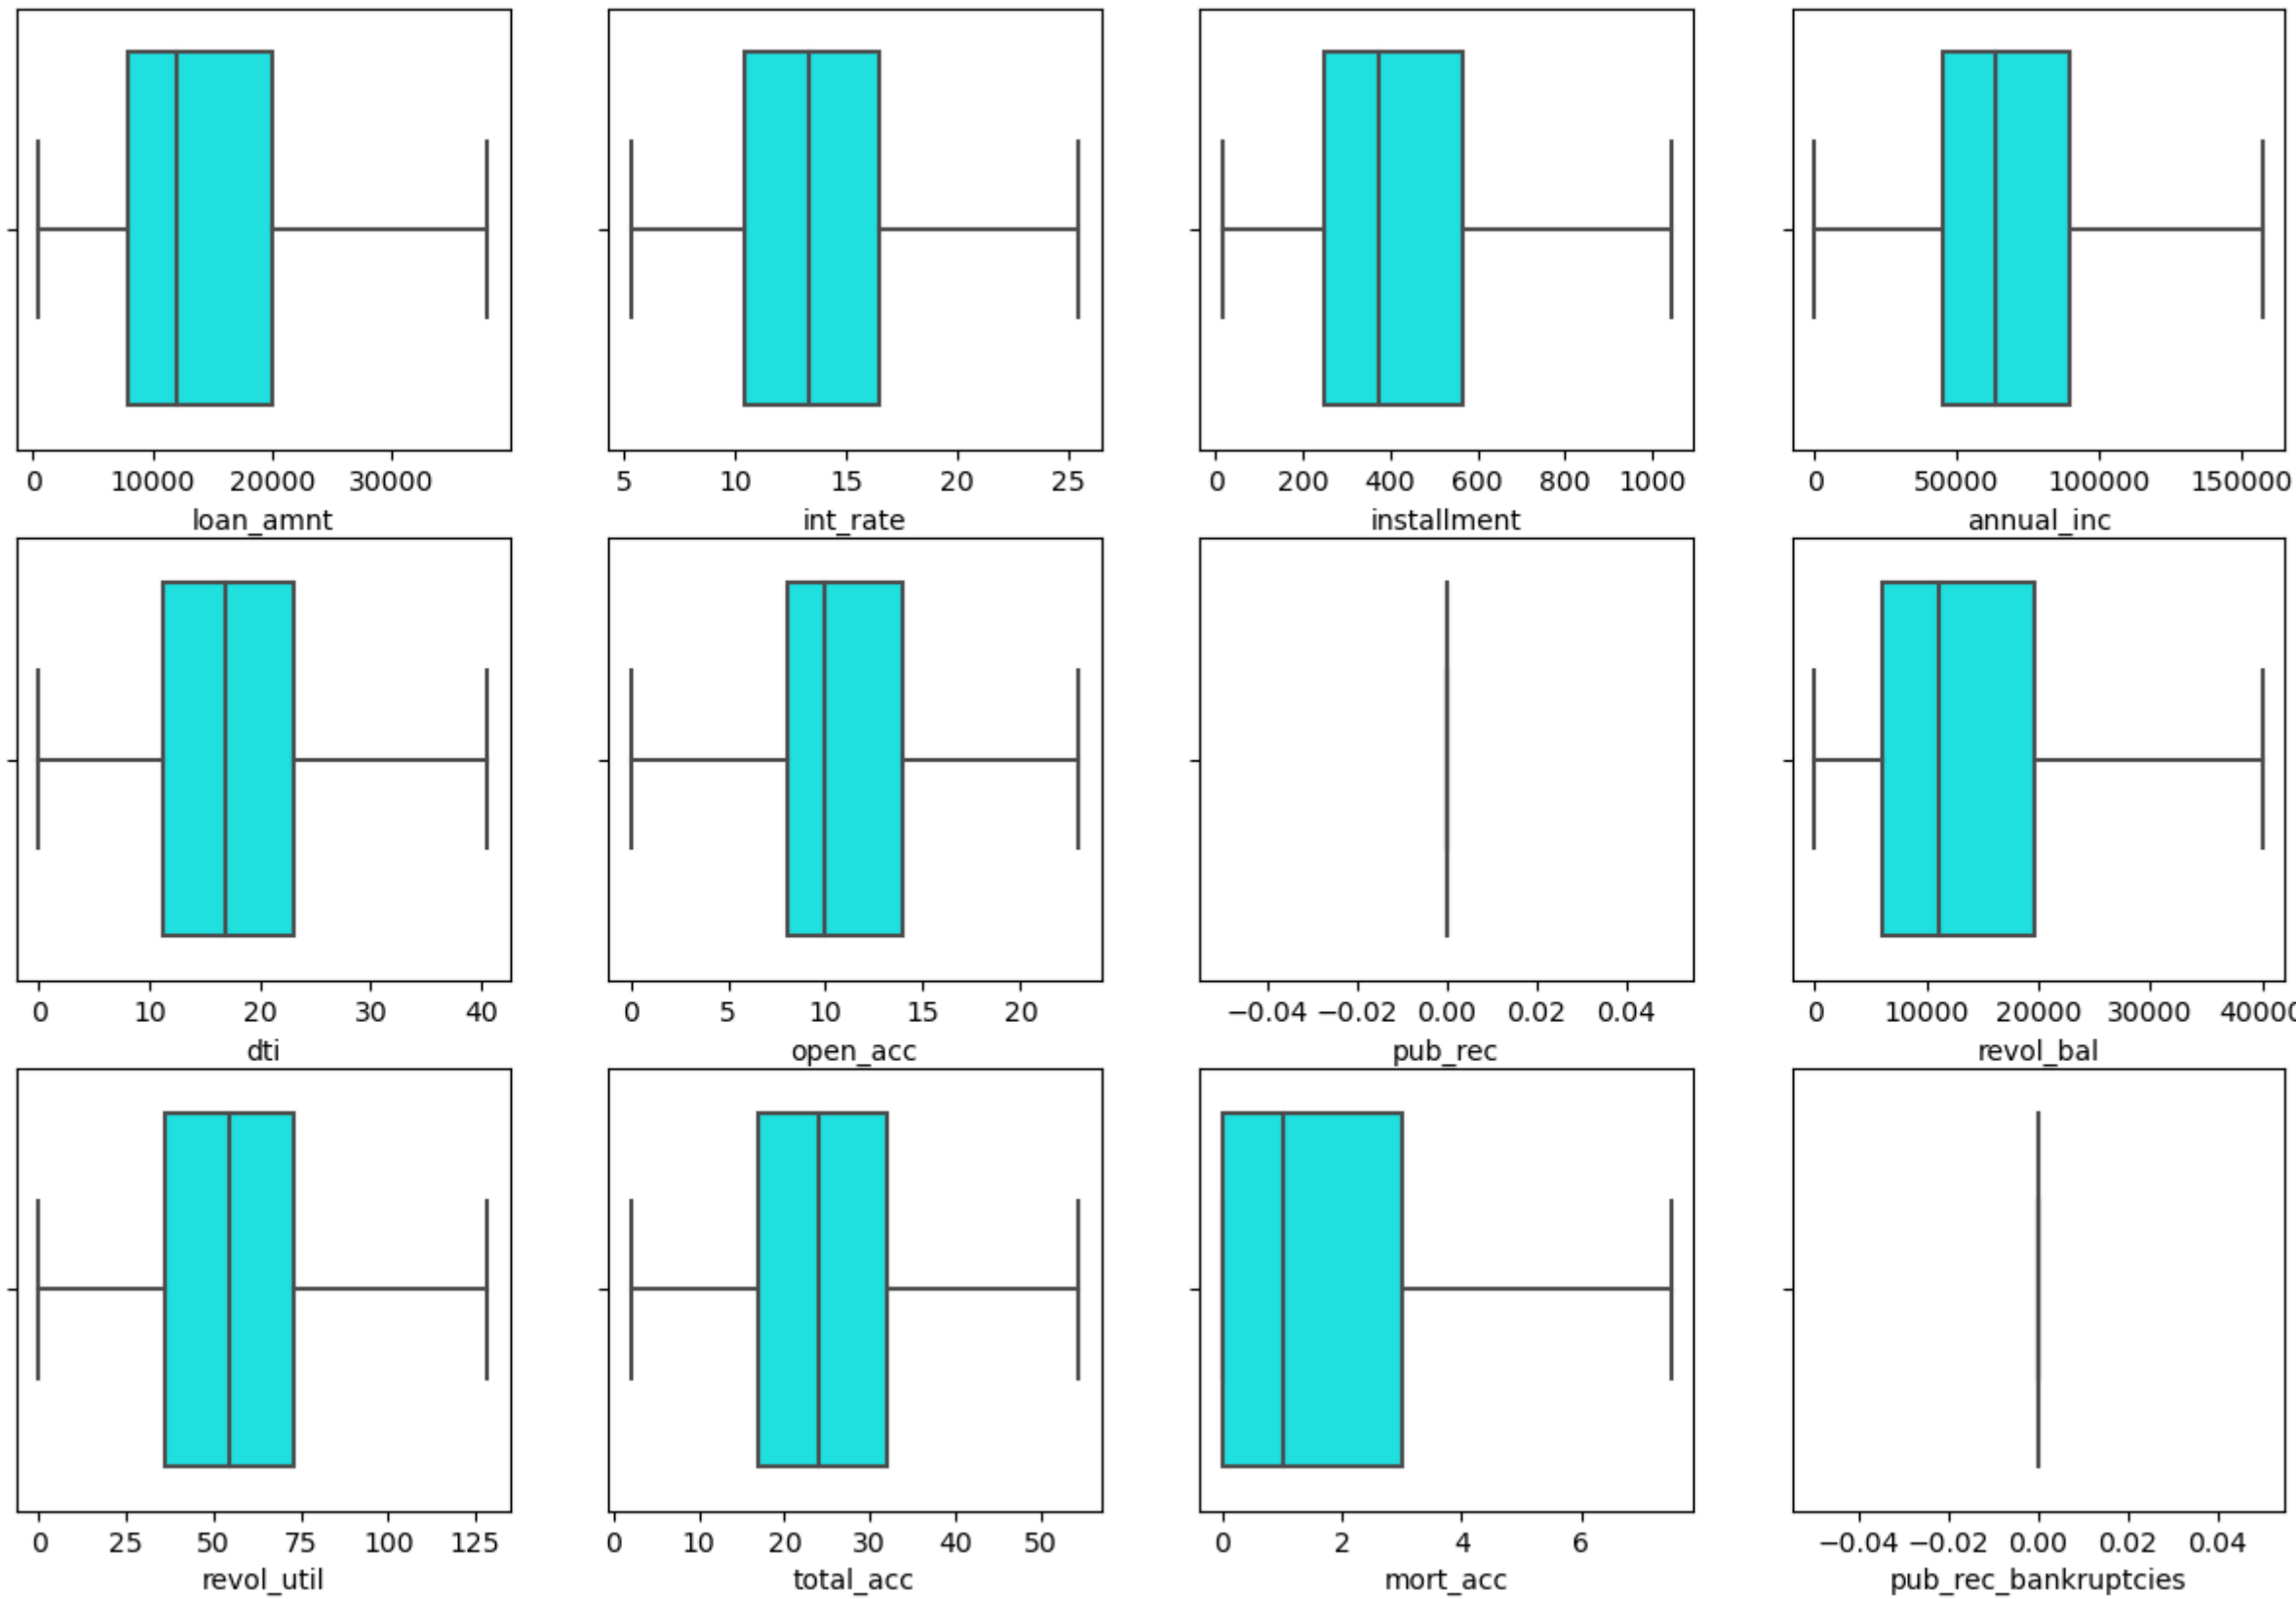
```
1 col = df2.columns[df2.dtypes == 'float64']
2 fig, ax = plt.subplots(3,4, figsize = (15,10))
3 c = 0
4 for i in range(3):
5     for j in range(4):
6         sns.boxplot(ax = ax[i,j], x = df2[col[c]], color = 'cyan')
7         c = c+1
```



In [65]:

```
1 # converting all outlier to upper limit
2 col = df2.columns[df2.dtypes == 'float64']
3
4 for i in col:
5     Q1=df2[i].quantile(0.25)
6     Q3=df2[i].quantile(0.75)
7     iqr = Q3-Q1
8     Upper_Whisker = Q3+1.5*iqr
9     df2.loc[df2[i] > Upper_Whisker,i] = Upper_Whisker
```

```
In [66]: 1 col = df2.columns[df2.dtypes == 'float64']
2 fig, ax = plt.subplots(3,4, figsize = (15,10))
3 c = 0
4 for i in range(3):
5     for j in range(4):
6         sns.boxplot(ax = ax[i,j], x = df2[col[c]], color = 'cyan')
7         c = c+1
```



Encoding of column from string to numeric

- home\_ownership
- verification\_status
- loan\_status
- purpose
- application\_type
- grade
- sub\_grade
- initial\_list\_status

```
In [67]: 1 df3 = df2.copy()
```

```
In [68]: 1 df3['loan_status'].unique()
```

Out[68]: array(['Fully Paid', 'Charged Off'], dtype=object)

```
In [69]: 1 df3['loan_status'] = df3['loan_status'].map({'Fully Paid':0, 'Charged Off':1})
```

```
In [70]: 1 # Target encoding of below 4 columns
2 cols= ['home_ownership','verification_status', 'purpose', 'application_type','grade','sub_grade','initial_list_status','emp_title', 'title']
3 for i in cols:
4     te = TargetEncoder()
5     df3[i] = df3[i].astype('category')
6     df3[i] = te.fit_transform(X = df3[i], y = df3['loan_status'])
```

Extracting only number from object

- term
- emp\_length

```
In [71]: 1 df3['term'] = df3['term'].map({' 36 months':36, ' 60 months': 60})
```

```
In [72]: 1 di = {'10+ years' : 10.5, '4 years' : 4, '< 1 year' : 0.5, '6 years': 6, '9 years' : 9, '2 years' : 2, '3 years' : 3, '8 years' : 8, '7 years' : 7, '5 years' : 5, '1 year' : 1}
2 df3['emp_length'] = df3['emp_length'].map(di)
3 df3['emp_length'].unique()
```

Out[72]: array([10.5, 4. , 0.5, 6. , 9. , 2. , 3. , 8. , 7. , 5. , 1. ])

Feature engineering for date column -> current date - given date in years

```
In [73]: 1 df3['issue_d'] = pd.to_datetime(df3['issue_d'], format='%b-%Y')
2 df3['issue_d_diff'] = date.today().year - df3['issue_d'].dt.year
```

```
In [74]: 1 df3['earliest_cr_line'] = pd.to_datetime(df3['earliest_cr_line'], format='%b-%Y')
2 df3['earliest_cr_line_diff'] = date.today().year - df3['earliest_cr_line'].dt.year
```

Drop columns which are not related to Target column

- issue\_d -> feature engineering done for this column
- earliest\_cr\_line -> feature engineering done for this column
- address -> wont impact target column

```
In [75]: 1 ncols = ['loan_amnt', 'term', 'int_rate', 'installment', 'grade', 'sub_grade',
2          'emp_length', 'home_ownership', 'annual_inc',
3          'verification_status', 'loan_status', 'purpose',
4          'dti', 'open_acc', 'pub_rec', 'revol_bal',
5          'revol_util', 'total_acc', 'initial_list_status', 'application_type',
6          'mort_acc', 'pub_rec_bankruptcies', 'issue_d_diff',
7          'earliest_cr_line_diff', 'emp_title', 'title']
8
9 df4 = df3[ncols]
```

Univarient analysis

```
In [76]: 1 col = df.columns
2         for i in col:
3             print('Unique values of',i, 'column =', df1[i].nunique())
4
5             print('-----')
```

Unique values of loan\_amnt column = 1397  
-----  
Unique values of term column = 2  
-----  
Unique values of int\_rate column = 566  
-----  
Unique values of installment column = 55706  
-----  
Unique values of grade column = 7  
-----  
Unique values of sub\_grade column = 35  
-----  
Unique values of emp\_title column = 173106  
-----  
Unique values of emp\_length column = 11  
-----  
Unique values of home\_ownership column = 6  
-----  
Unique values of annual\_inc column = 27197  
-----  
Unique values of verification\_status column = 3  
-----  
Unique values of issue\_d column = 115  
-----  
Unique values of loan\_status column = 2  
-----  
Unique values of purpose column = 14  
-----  
Unique values of title column = 48818  
-----  
Unique values of dti column = 4262  
-----  
Unique values of earliest\_cr\_line column = 684  
-----  
Unique values of open\_acc column = 61  
-----  
Unique values of pub\_rec column = 20  
-----  
Unique values of revol\_bal column = 55622  
-----  
Unique values of revol\_util column = 1227  
-----  
Unique values of total\_acc column = 118  
-----  
Unique values of initial\_list\_status column = 2  
-----  
Unique values of application\_type column = 3  
-----  
Unique values of mort\_acc column = 34  
-----  
Unique values of pub\_rec\_bankruptcies column = 10  
-----  
Unique values of address column = 393700  
-----

Unique value counts

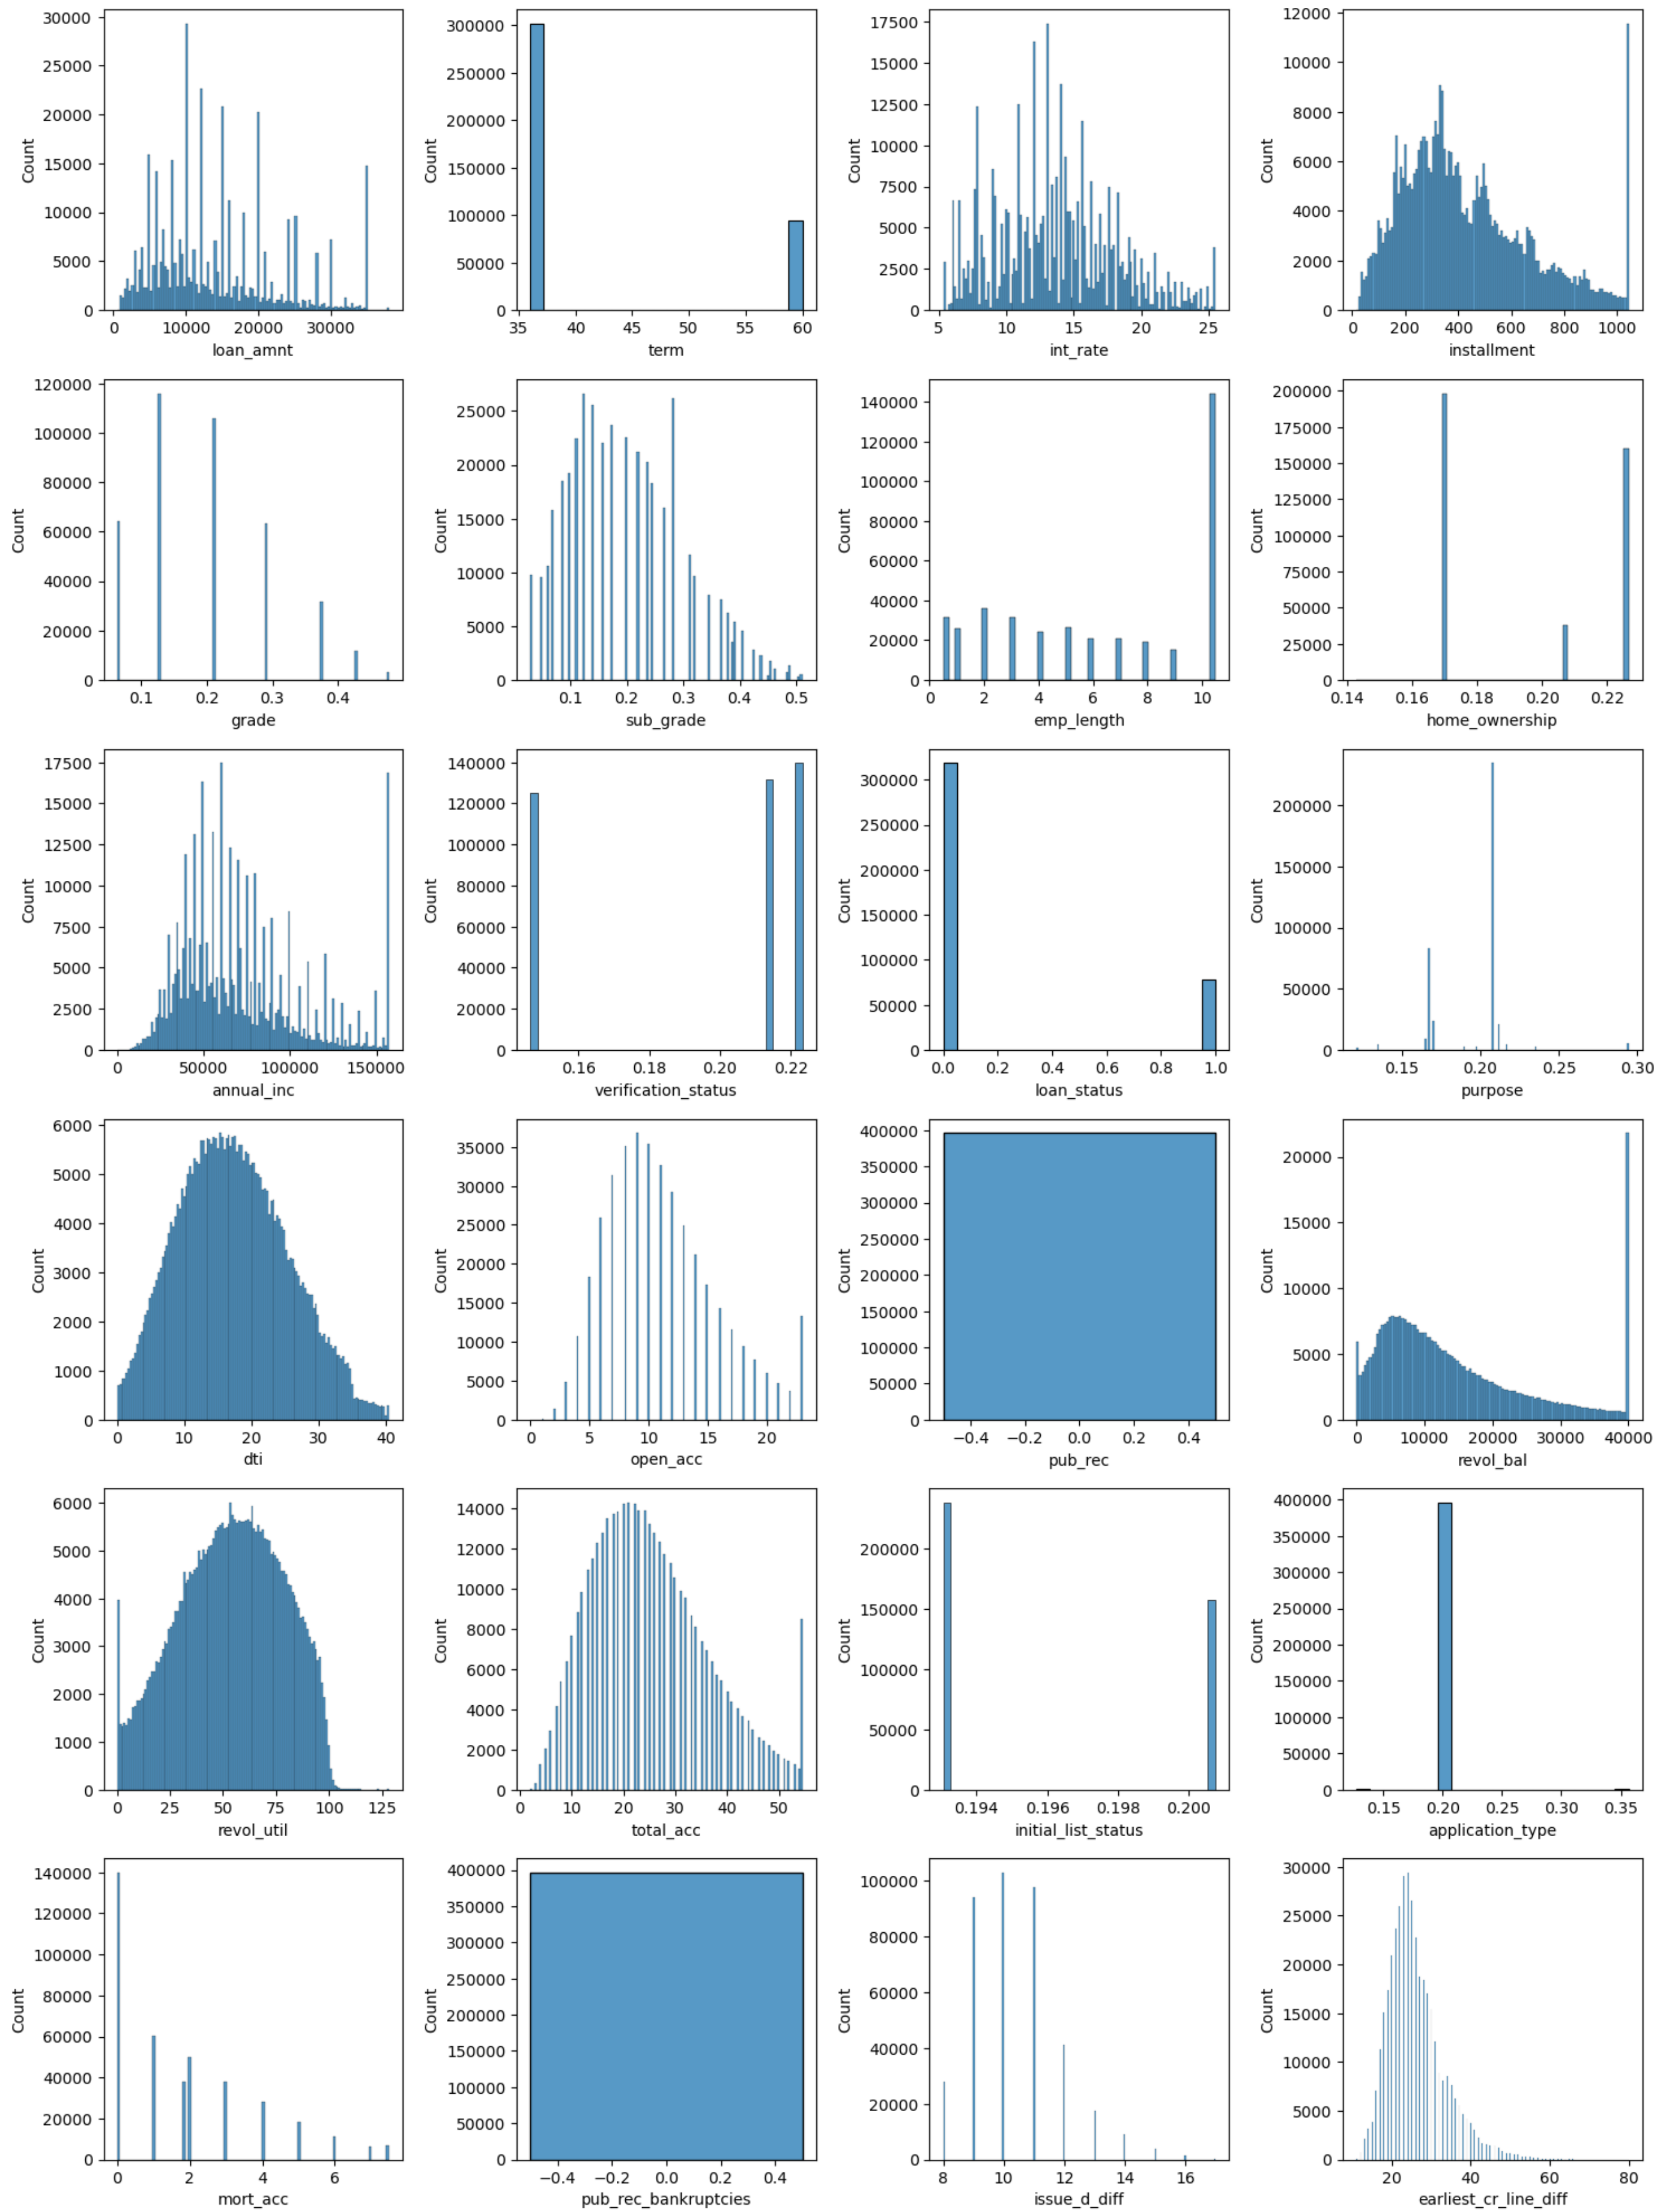
```
In [77]: 1 col = df.columns
2         for i in col:
3             print('Unique values of',i, 'column')
4             print((df1[i].value_counts().head(3)/df.shape[0]) * 100)
5             print('-----')
```

10.0 8.949070  
8.0 8.872308  
Name: open\_acc, dtype: float64  
-----  
Unique values of pub\_rec column  
0.0 85.415751  
1.0 12.559402  
2.0 1.382724  
Name: pub\_rec, dtype: float64  
-----  
Unique values of revol\_bal column  
0.0 0.537333  
5655.0 0.010353  
6095.0 0.009595  
Name: revol\_bal, dtype: float64  
-----  
Unique values of revol\_util column  
0.0 0.558796  
53.0 0.189885  
60.0 0.186602



Histogram

```
In [80]: 1 col = df4.columns
2 fig, ax = plt.subplots(6,4, figsize = (15,20))
3 c = 0
4 for i in range(6):
5     for j in range(4):
6         sns.histplot(ax = ax[i,j], x = df4[col[c]])
7         c = c+1
8 plt.tight_layout()
9 plt.show()
```



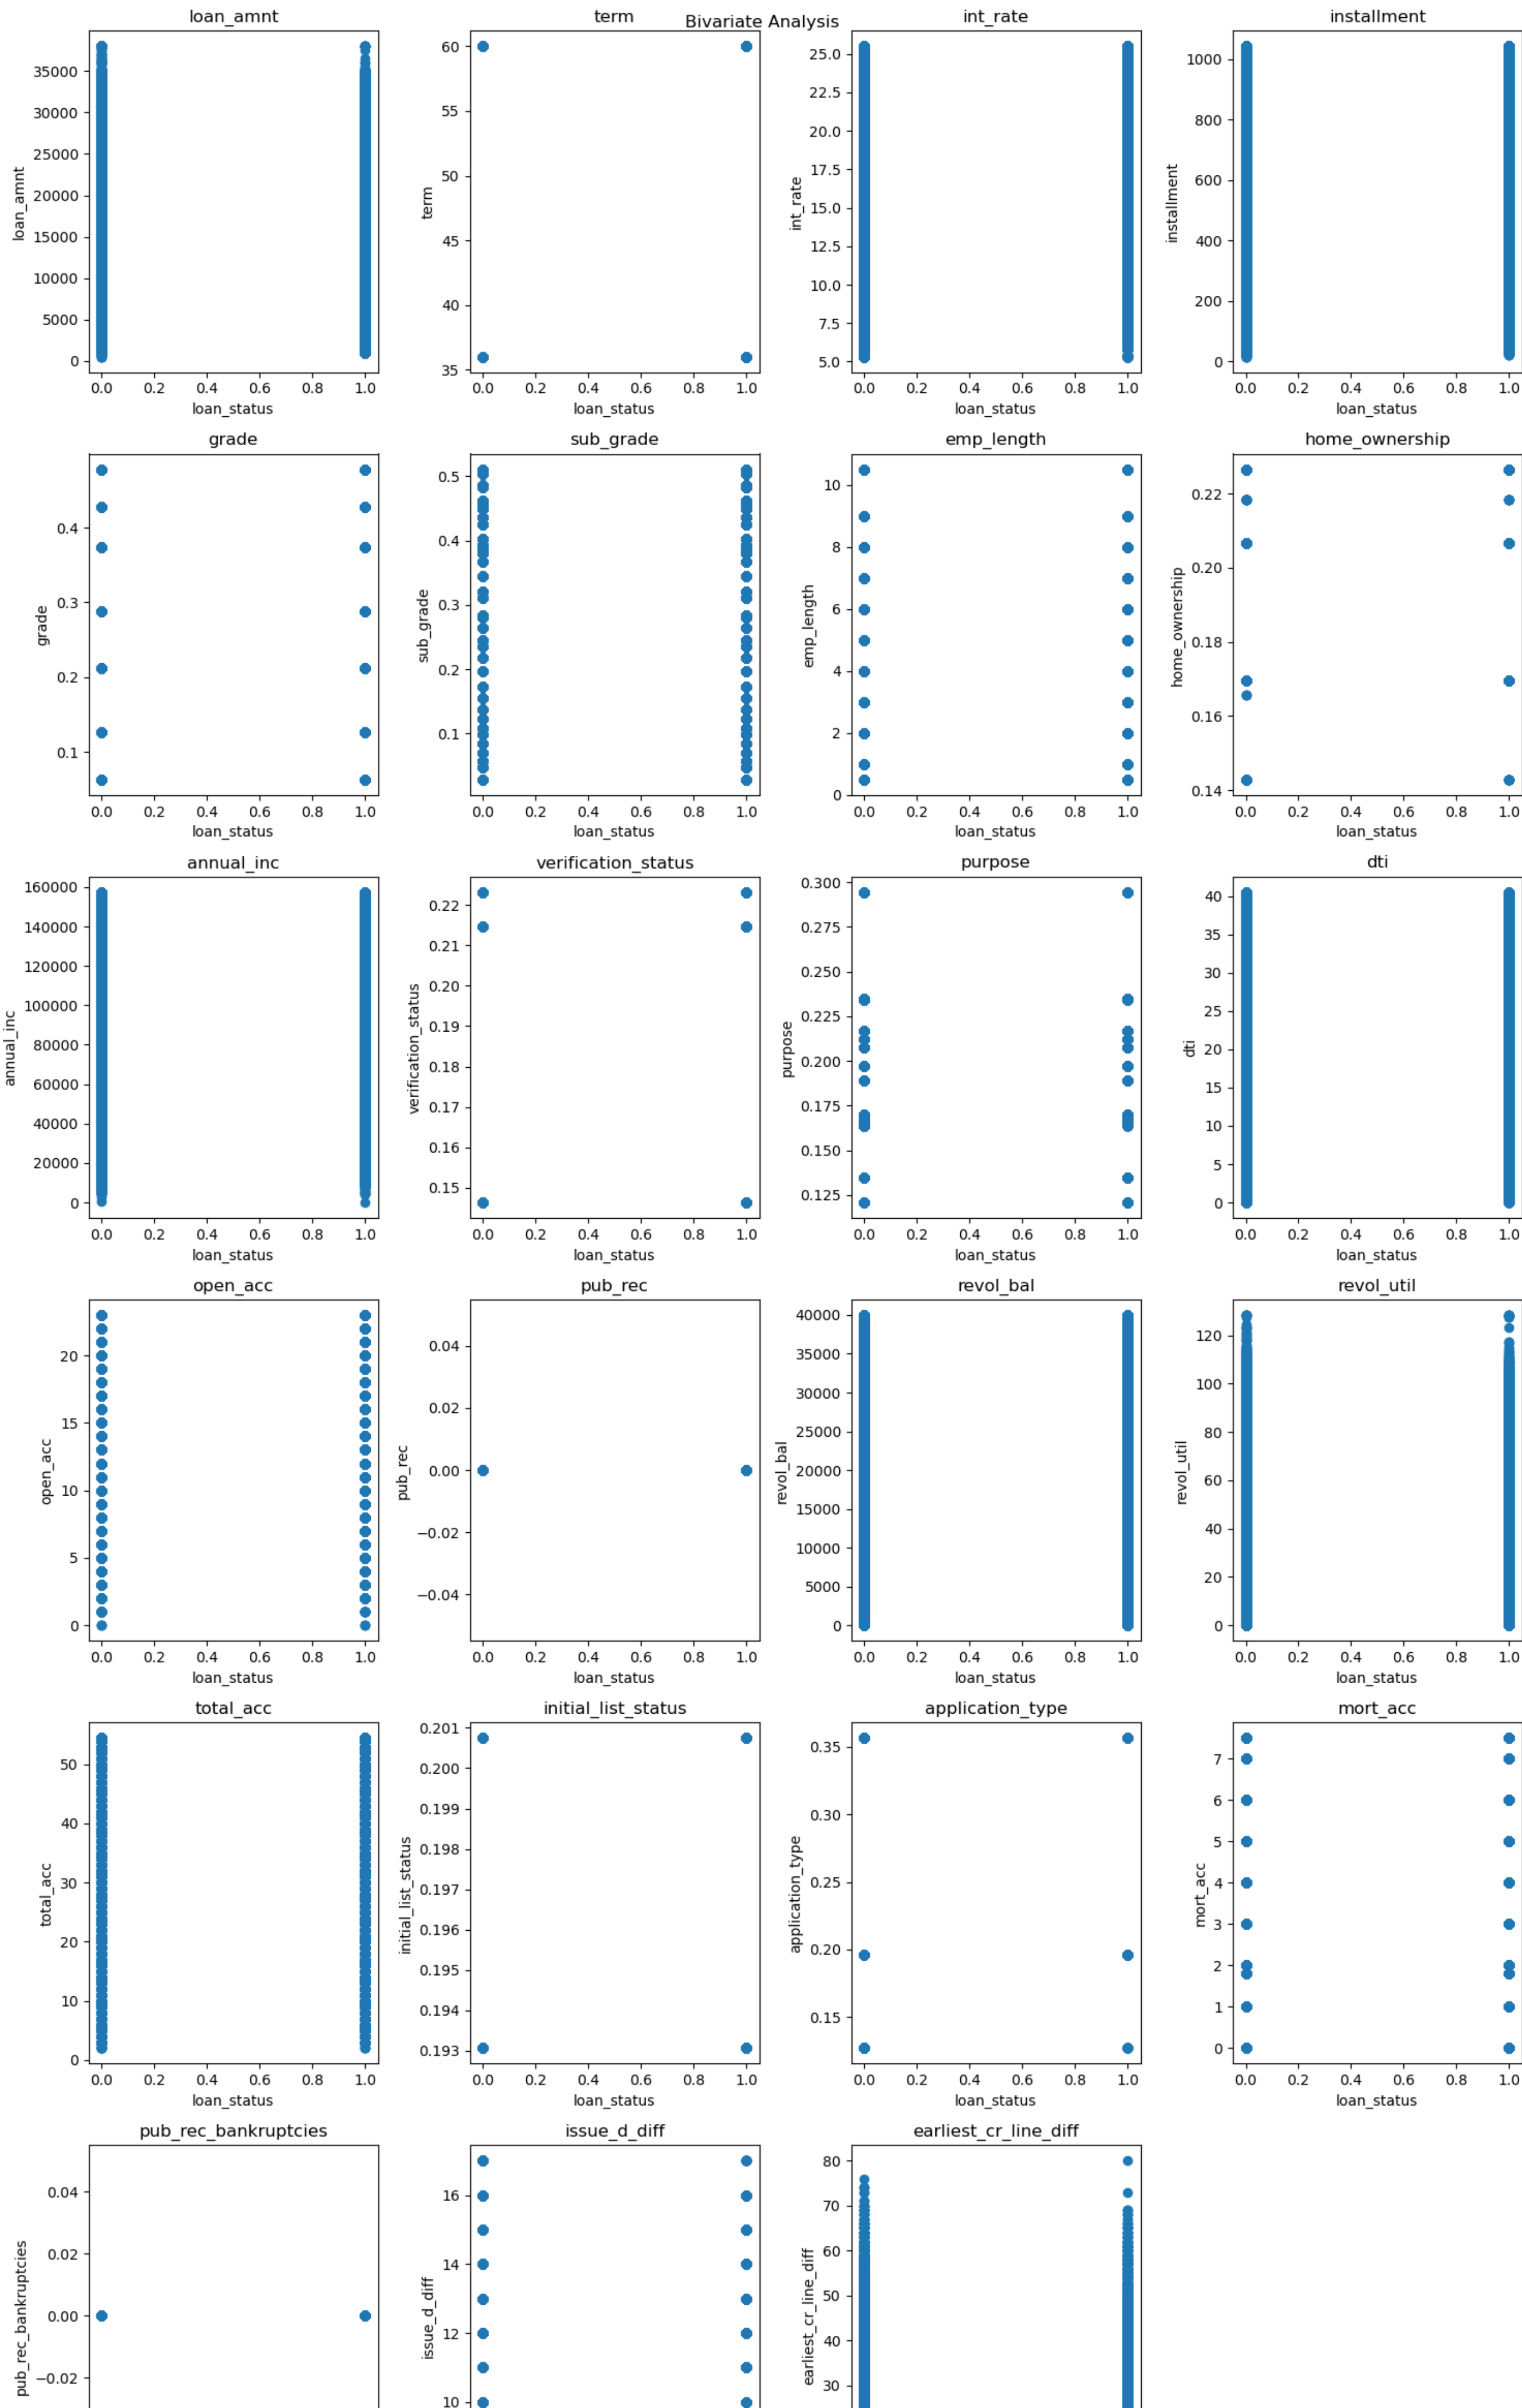
## Bivariate Analysis


```
In [82]: 1 col = df4.columns.drop(['loan_status'])
2 fig, ax = plt.subplots(6,4, figsize = (15,25))
3 fig.suptitle('Bivariate Analysis')
4 #ax[2,1].set_axis_off()
5 ax[5,3].set_axis_off()
6 c = 0
7 print(len(col))
8 for i in range(6):
9     for j in range(4):
10         #print( df[col[c]])
11         if c == 23:
12             break
13         ax[i,j].scatter(x = df4['loan_status'], y = df4[col[c]])
14         ax[i,j].set_xlabel('loan_status')
15         ax[i,j].set_ylabel(col[c])
16         ax[i,j].set_title(col[c])
17         c = c +1
18 plt.tight_layout()
19 plt.show()
```

25







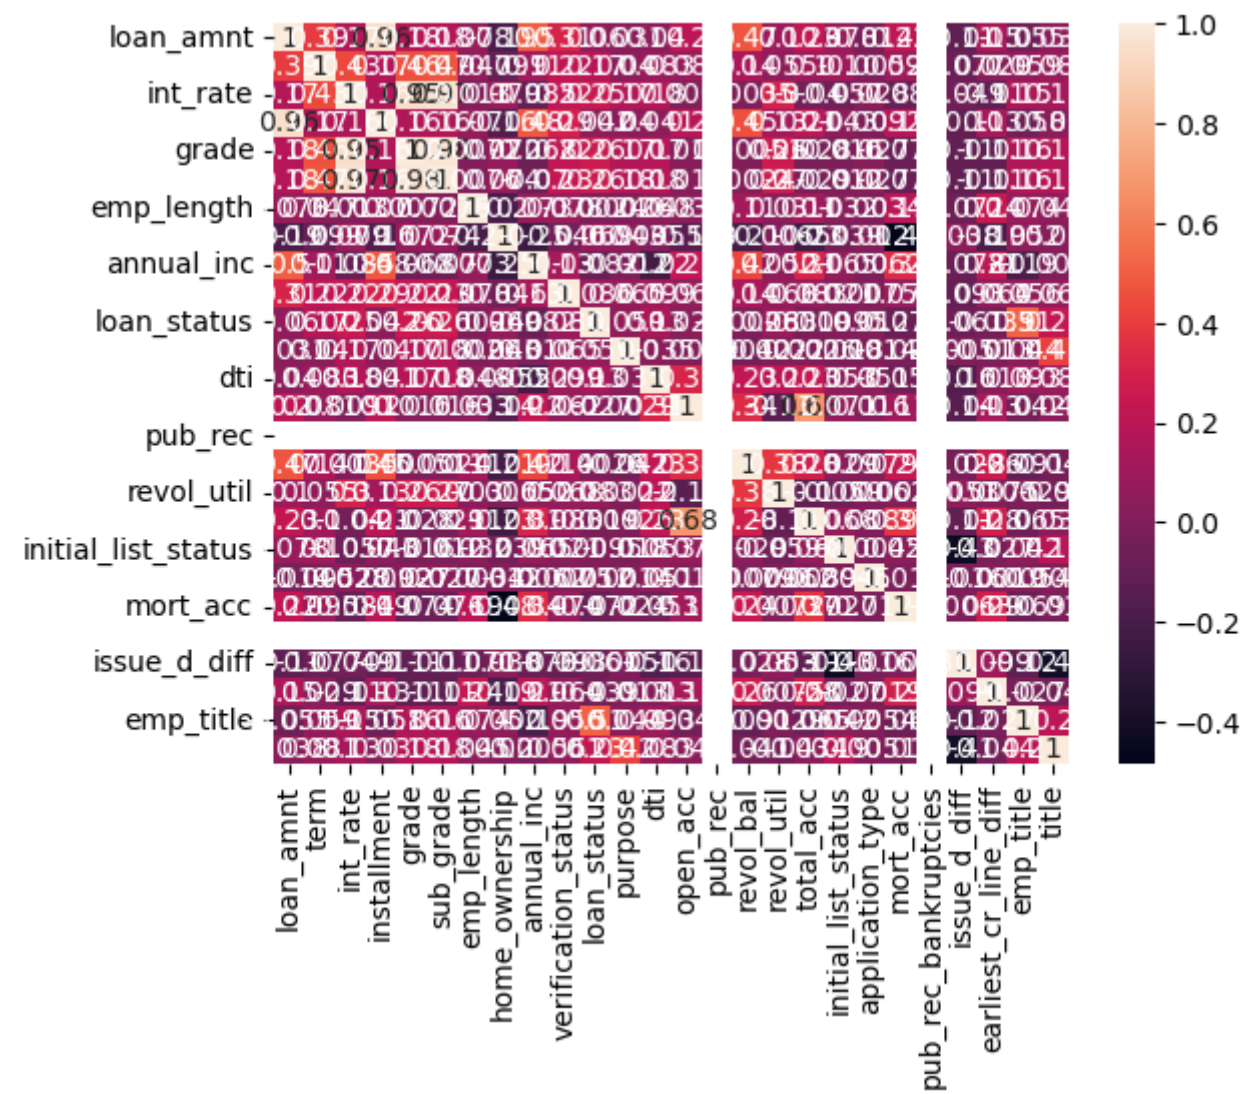


Bivariate insights :

- Mort\_acc, total\_acc, revol\_bal as correlation with Loan status column

```
In [78]: 1 sns.heatmap(df4.corr(method='pearson'), annot=True)
```

Out[78]: <Axes: >



```
In [79]: 1 df4.corr(method='pearson').loc['loan_status'].sort_values(ascending = False)
```

Out[79]: loan\_status 1.000000  
emp\_title 0.508090  
sub\_grade 0.263801  
grade 0.257886  
int\_rate 0.248077  
title 0.225088  
term 0.173246  
dti 0.132507  
verification\_status 0.085618  
revol\_util 0.082505  
home\_ownership 0.068534  
loan\_amnt 0.059898  
purpose 0.059394  
installment 0.042407  
open\_acc 0.027475  
application\_type 0.012268  
initial\_list\_status 0.009489  
emp\_length -0.002394  
revol\_bal -0.002575  
total\_acc -0.018826  
earliest\_cr\_line\_diff -0.038928  
issue\_d\_diff -0.060502  
mort\_acc -0.072250  
annual\_inc -0.082195  
pub\_rec NaN  
pub\_rec\_bankruptcies NaN  
Name: loan\_status, dtype: float64

Heatmap insights :

- Target variable 'Loan status' as highest positive correlation with sub\_grade and lowest negative correlation in mort\_acc


Data visualization

- Data doesnt have any general trend with lot of misclassification.

```
In [93]: 1 # feature reduction to 2 feature by PCA  
2 pca = PCA(n_components = 2)  
3 v = pca.fit_transform(df4)  
4 y = df4['loan_status']
```

```
In [95]: 1 sns.scatterplot(data = v, x = v[:,0], y = v[:,1], hue = y)
```

Out[95]: <Axes: >




### Inbalance check

- Target column is imbalance with 80 : 20 ratio. Class 0 : 80% data and Class 1 : 20% data
- With existing condition if we build model then recall score will be low. To improve recall data should be balanced.

```
In [87]: 1 sns.countplot(data = df4, x = 'loan_status')
```

Out[87]: <Axes: xlabel='loan\_status', ylabel='count'>



```
In [88]: 1 df4['loan_status'].value_counts(normalize = True)*100
```

Out[88]: 0 80.387092  
1 19.612908  
Name: loan\_status, dtype: float64

```
In [164]: 1 df['loan_status'].value_counts()
```

Out[164]: Fully Paid 318357  
Charged Off 77673  
Name: loan\_status, dtype: int64

### Split data into Train and test

- Data splited in train and test data. Train is used for assumption check, balancing data, building model and tuning for best lambda and regularization.

```
In [99]: 1 X = df4.drop(columns = ['loan_status'])  
2 y = df4['loan_status']
```

```
In [157]: 1 X_train_val, X_test, y_train_val, y_test = train_test_split(X, y, test_size = 0.2, random_state = 1)  
2 print('Train shape',X_train_val.shape, y_train_val.shape)  
3 print('Test shape',X_test.shape, y_test.shape)
```

Train shape (316824, 25) (316824,)  
Test shape (79206, 25) (79206,)

### Logistic regression Assumption checking

```
In [103]: 1 col_names = X_train_val.columns  
2  
3 # scale the model for checking assumption  
4 scale = StandardScaler()  
5 data_scale = scale.fit_transform(X_train_val)  
6 X_asm_sc = pd.DataFrame(data_scale, columns = col_names)  
7 y_asm = y
```

### Multicollinearity check by VIF score

- As per 1st multicollinearity check, 5 columns as high multicollinearity VIF values.
- With removal of 'installment', 'sub\_grade','grade' these columns, VIF value got reduced below 5.



```
In [104]: 1 # create dataframe to VIF values
2 vif = pd.DataFrame()
3
4 # create features column for comparision
5 vif['Features'] = X_asm_sc.columns
6
7 # create VIF values for all independent columns
8 vif['Vif values'] = [variance_inflation_factor(X_asm_sc.values, i) for i in range(X_asm_sc.shape[1])]
9
10 # round values
11 vif['Vif values'] = round(vif['Vif values'],2)
12
13 # sort by values in decending order
14 vif= vif.sort_values(by ='Vif values', ascending = False )
15 vif
```

C:\Users\trtej\anaconda3\lib\site-packages\statsmodels\regression\linear\_model.py:1754: RuntimeWarning: invalid value encountered in double\_scalars  
return 1 - self.ssr/self.uncentered\_tss

Out[104]:

	Features	Vif values
0	loan_amnt	57.01
3	installment	49.27
5	sub_grade	40.33
4	grade	22.59
2	int_rate	19.51
1	term	5.82
16	total_acc	2.38
12	open_acc	2.30
14	revol_bal	2.05
8	annual_inc	1.88
19	mort_acc	1.62
21	issue_d_diff	1.60
24	title	1.60
15	revol_util	1.55
11	dti	1.49
7	home_ownership	1.36
10	purpose	1.30
17	initial_list_status	1.27
22	earliest_cr_line_diff	1.26
9	verification_status	1.16
6	emp_length	1.12
23	emp_title	1.12
18	application_type	1.00
13	pub_rec	NaN
20	pub_rec_bankruptcies	NaN

```
In [105]: 1 ## 2nd Multicollinearity check by removing installment from data
2 # create dataframe to VIF values
3 vif = pd.DataFrame()
4 X_asm_sc_2 = X_asm_sc.drop(columns = ['installment'])
5 # create features column for comparision
6 vif['Features'] = X_asm_sc_2.columns
7
8 # create VIF values for all independent columns
9 vif['Vif values'] = [variance_inflation_factor(X_asm_sc_2.values, i) for i in range(X_asm_sc_2.shape[1])]
10
11 # round values
12 vif['Vif values'] = round(vif['Vif values'],2)
13
14 # sort by values in decending order
15 vif= vif.sort_values(by ='Vif values', ascending = False )
16 vif
```

C:\Users\trtej\anaconda3\lib\site-packages\statsmodels\regression\linear\_model.py:1754: RuntimeWarning: invalid value encountered in double\_scalars  
return 1 - self.ssr/self.uncentered\_tss

Out[105]:

	Features	Vif values
4	sub_grade	40.33
3	grade	22.59
2	int_rate	18.96
15	total_acc	2.38
11	open_acc	2.29
13	revol_bal	2.05
7	annual_inc	1.88
0	loan_amnt	1.86
18	mort_acc	1.62
20	issue_d_diff	1.60
23	title	1.60
14	revol_util	1.55
1	term	1.54
10	dti	1.49
6	home_ownership	1.36
9	purpose	1.30
16	initial_list_status	1.27
21	earliest_cr_line_diff	1.26
8	verification_status	1.16
5	emp_length	1.12
22	emp_title	1.12
17	application_type	1.00
12	pub_rec	NaN
19	pub_rec_bankruptcies	NaN

```
In [106]: 1 ## 3rd Multicollinearity check by removing installment,sub_grade from data
2 # create dataframe to VIF values
3 vif = pd.DataFrame()
4 X_asm_sc_3 = X_asm_sc.drop(columns = ['installment', 'sub_grade'])
5 # create features column for comparision
6 vif['Features'] = X_asm_sc_3.columns
7
8 # create VIF values for all independent columns
9 vif['Vif values'] = [variance_inflation_factor(X_asm_sc_3.values, i) for i in range(X_asm_sc_3.shape[1])]
10
11 # round values
12 vif['Vif values'] = round(vif['Vif values'],2)
13
14 # sort by values in decending order
15 vif= vif.sort_values(by ='Vif values', ascending = False )
16 vif
```

C:\Users\trtej\anaconda3\lib\site-packages\statsmodels\regression\linear\_model.py:1754: RuntimeWarning: invalid value encountered in double\_scalars  
return 1 - self.ssr/self.uncentered\_tss

Out[106]:

	Features	Vif values
2	int_rate	10.97
3	grade	10.90
14	total_acc	2.38
10	open_acc	2.29
12	revol_bal	2.05
6	annual_inc	1.88
0	loan_amnt	1.86
17	mort_acc	1.62
19	issue_d_diff	1.59
22	title	1.58
13	revol_util	1.55
1	term	1.51
9	dti	1.49
5	home_ownership	1.36
8	purpose	1.30
15	initial_list_status	1.26
20	earliest_cr_line_diff	1.26
7	verification_status	1.16
4	emp_length	1.12
21	emp_title	1.12
16	application_type	1.00
11	pub_rec	NaN
18	pub_rec_bankruptcies	NaN

```
In [107]: 1 ## 4th Multicollinearity check by removing installment,sub_grade, grade from data
2 # create dataframe to VIF values
3 vif = pd.DataFrame()
4 X_asm_sc_4 = X_asm_sc.drop(columns = ['installment', 'sub_grade','grade'])
5 # create features column for comparision
6 vif['Features'] = X_asm_sc_4.columns
7
8 # create VIF values for all independent columns
9 vif['Vif values'] = [variance_inflation_factor(X_asm_sc_4.values, i) for i in range(X_asm_sc_4.shape[1])]
10
11 # round values
12 vif['Vif values'] = round(vif['Vif values'],2)
13
14 # sort by values in decending order
15 vif= vif.sort_values(by ='Vif values', ascending = False )
16 vif
```

C:\Users\trtej\anaconda3\lib\site-packages\statsmodels\regression\linear\_model.py:1754: RuntimeWarning: invalid value encountered in double\_scalars  
return 1 - self.ssr/self.uncentered\_tss

Out[107]:

	Features	Vif values
13	total_acc	2.38
9	open_acc	2.29
11	revol_bal	2.05
5	annual_inc	1.88
0	loan_amnt	1.86
16	mort_acc	1.62
2	int_rate	1.61
18	issue_d_diff	1.57
21	title	1.57
12	revol_util	1.55
8	dti	1.49
1	term	1.48
4	home_ownership	1.36
7	purpose	1.30
14	initial_list_status	1.26
19	earliest_cr_line_diff	1.26
6	verification_status	1.16
3	emp_length	1.12
20	emp_title	1.12
15	application_type	1.00
10	pub_rec	NaN
17	pub_rec_bankruptcies	NaN

```
In [109]: 1 X_train_val= X_train_val.drop(columns = ['installment', 'sub_grade','grade'])
```

balance data using SMOTE technique

- New points are created using KNN technique to balance data so recall and precision can be balanced.



```
In [111]: 1 sm = SMOTE(random_state=42)
          2 X_res, y_res = sm.fit_resample(X_train_val, y_train_val)

In [120]: 1 X_train_val.shape

Out[120]: (316824, 22)

In [118]: 1 X_res.shape

Out[118]: (509602, 22)

In [119]: 1 y_res.shape

Out[119]: (509602,)
```

-----

Split data - Train and validation

- split train\_val data to train and validation data.
- Train data to build model
- Validation data to hyper parameter tuning.

```
In [150]: 1 X_train, X_val, y_train, y_val = train_test_split(X_res, y_res, test_size = 0.25, random_state = 1)
          2 print('Train shape',X_train.shape, y_train.shape)
          3 print('Val shape',X_val.shape,y_val.shape)
          4

Train shape (382201, 22) (382201,)
Val shape (127401, 22) (127401,)
```

Scaling data for Train and validation data

```
In [151]: 1 # Scale the model
          2 scale = StandardScaler()
          3 X_train = scale.fit_transform(X_train)
          4 X_val = scale.transform(X_val)
          5
          6 y_train = y_train.values
          7 y_val = y_val.values

In [124]: 1 np.unique(y_train,return_counts=True)

Out[124]: (array([0, 1], dtype=int64), array([191050, 191151], dtype=int64))
```

Logistic regression by Sklearn

- Train and validation data:
  - Model build on train data.
  - F1 score - 0.82 and AU PR curve is 0.86 (after hyperparameter tuning - Lamdba, Regularizaiont).
  -


- Test data:
  - F1 score is 0.62 and Area under PR curve is 0.67.
  - Due to data hetrogenous and lot misclassification. F1 score is reduced.

```
In [244]: 1 # Define Logistic regression
          2 model = LogisticRegression()
          3 model.fit(X_train, y_train)
          4
          5 # accuracy on train
          6 accuracy = model.score(X_train, y_train)
          7
          8 # accuracy on val data
          9 acc_val = model.score(X_val, y_val)
          10
          11 # summary
          12 print('accuracy on train', accuracy)
          13 print('accuracy on val', acc_val)
          14
          15 # predict y
          16 y_train_pred = model.predict(X_train)
          17 y_val_pred = model.predict(X_val)

accuracy on train 0.823168960834744
accuracy on val 0.8218852285303883
```

```
In [126]: 1 #confusion matrix on validation data
2
3 cmp = confusion_matrix(y_val,y_val_pred )
4 disp = ConfusionMatrixDisplay(cmp)
5 disp.plot()
```

Out[126]: <sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay at 0x1f113246530>




```
In [127]: 1 # precision and recall
2 print('Precision score on valdata ',precision_score(y_val, y_val_pred))
3 print('Recall score on val data ',recall_score(y_val, y_val_pred))
```

Precision score on valdata 0.8197146157929247  
Recall score on val data 0.8249175176747839

```
In [128]: 1 # f1score
2 print('F1_score', f1_score(y_val, y_val_pred))
```


F1\_score 0.8223078368727683

```
In [129]: 1 # check ROC curve and AUROC value
2 fpr, tpr, th = roc_curve(y_val, y_val_pred)
3 plt.plot(fpr, tpr)
4 plt.plot(fpr,fpr, '--',color = 'red')
5 plt.title('ROC curve')
6 plt.xlabel('FPR')
7 plt.ylabel('TPR')
8 plt.show()
9
10 ## AUROC score
11 roc_auc_score(y_val, y_val_pred)
```



Out[129]: 0.8218876305413652

```
In [130]: 1 # Check PR curve and AU PR curve value
2 pr, rc, th = precision_recall_curve(y_val, y_val_pred)
3 plt.plot(rc,pr)
4 plt.title('Pr curve')
5 plt.xlabel('Recall')
6 plt.ylabel('Precision')
7 plt.show()
8
9 # AUC score of PR curve
10 print('AUC score of PR curve :',auc(rc, pr))
```



AUC score of PR curve : 0.8660519871740393

### Hyperparameter tuning to be done to get best f1 score

```
In [133]: 1 ## grid search CV
2 from sklearn.model_selection import GridSearchCV
3 from sklearn.metrics import f1_score, make_scorer
4
5 param_grid = {'C' : [0.001,0.01,1,5,10,15,20], 'penalty' : ['l2','l1']}
6
7 grid = GridSearchCV(estimator = LogisticRegression(), param_grid = param_grid, scoring = make_scorer(f1_score))
8 grid.fit(X_train, y_train)
```

C:\Users\trtej\anaconda3\lib\site-packages\sklearn\model\_selection\\_validation.py:378: FitFailedWarning: 35 fits failed out of a total of 70. The score on these train-test partitions for these parameters will be set to nan. If these failures are not expected, you can try to debug them by setting error\_score='raise'.

Below are more details about the failures:

-----

35 fits failed with the following error:

Traceback (most recent call last):

File "C:\Users\trtej\anaconda3\lib\site-packages\sklearn\model\_selection\\_validation.py", line 686, in \_fit\_and\_score

estimator.fit(X\_train, y\_train, \*\*fit\_params)

File "C:\Users\trtej\anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py", line 1162, in fit

solver = \_check\_solver(self.solver, self.penalty, self.dual)

File "C:\Users\trtej\anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py", line 54, in \_check\_solver

raise ValueError(

ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

warnings.warn(some\_fits\_failed\_message, FitFailedWarning)

C:\Users\trtej\anaconda3\lib\site-packages\sklearn\model\_selection\\_search.py:952: UserWarning: One or more of the test scores are non-finite: [0.82373168 nan 0.82378161 nan 0.82376504 nan nan 0.82376197 nan 0.82376197 nan 0.82376504 nan 0.82376197 nan]

warnings.warn(

Out[133]:

GridSearchCV

estimator: LogisticRegression

LogisticRegression

```
In [134]: 1 grid.best_params_
```


Out[134]: {'C': 0.01, 'penalty': 'l2'}

```
In [135]: 1 # Define Logistic regression
2 model1 = LogisticRegression(penalty='l2',C=0.1)
3 model1.fit(X_train, y_train)
4
5 # accuracy on train
6 accuracy = model1.score(X_train, y_train)
7
8 # accuracy on val data
9 acc_val = model1.score(X_val, y_val)
10
11 # summary
12 print('accuracy on train', accuracy)
13 print('accuracy on val', acc_val)
14
15 # predict y
16 y_train_pred1 = model1.predict(X_train)
17 y_val_pred1 = model1.predict(X_val)
```

accuracy on train 0.8231637279860597  
accuracy on val 0.8218773792984356

```
In [136]: 1 #confusion matrix on validation data
2
3 cmp = confusion_matrix(y_val,y_val_pred1 )
4 disp = ConfusionMatrixDisplay(cmp)
5 disp.plot()
```


Out[136]: <sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay at 0x1f15b3ceda0>



```
In [137]: 1 # precision and recall
2 print('Precision score on valdata ',precision_score(y_val, y_val_pred1))
3 print('Recall score on val data',recall_score(y_val, y_val_pred1))
4
5 # f1score
6 print('F1_score', f1_score(y_val, y_val_pred1))
```


Precision score on valdata 0.8197118011646605  
Recall score on val data 0.8249018067556952  
F1\_score 0.8222986147545476

```
In [138]: 1 # check ROC curve and AUROC value
2 fpr, tpr, th = roc_curve(y_val, y_val_pred1)
3 plt.plot(fpr, tpr)
4 plt.plot(fpr,fpr, '--',color = 'red')
5 plt.title('ROC curve')
6 plt.xlabel('FPR')
7 plt.ylabel('TPR')
8 plt.show()
9
10 ## AUROC score
11 roc_auc_score(y_val, y_val_pred)
```



Out[138]: 0.8218876305413652

```
In [139]: 1 # Check PR curve and AU PR curve value
2 pr, rc, th = precision_recall_curve(y_val, y_val_pred1)
3 plt.plot(rc,pr)
4 plt.title('Pr curve')
5 plt.xlabel('Recall')
6 plt.ylabel('Precision')
7 plt.show()
8
9 # AUC score of PR curve
10 print('AUC score of PR curve :',auc(rc, pr))
```




AUC score of PR curve : 0.8660466490163391

```
In [158]: 1 # check on test data
2 X_test= X_test.drop(columns = ['installment', 'sub_grade','grade'])
3 X_test = scale.transform(X_test)
```

```
In [159]: 1 y_test_pred1 = model1.predict(X_test)
```

```
In [160]: 1 #confusion matrix on validation data
2
3 cmp = confusion_matrix(y_test,y_test_pred1 )
4 disp = ConfusionMatrixDisplay(cmp)
5 disp.plot()
```


Out[160]: <sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay at 0x1f15de74c10>



```
In [161]: 1 # precision and recall
2 print('Precision score on valdata ',precision_score(y_test,y_test_pred1))
3 print('Recall score on val data',recall_score(y_test,y_test_pred1))
4
5 # f1score
6 print('F1_score', f1_score(y_test,y_test_pred1))
```


Precision score on valdata 0.5169338970988404  
Recall score on val data 0.7890095846645367  
F1\_score 0.6246300933302982

```
In [243]: 1 # check ROC curve and AUROC value
2 fpr, tpr, th = roc_curve(y_test,y_test_pred1)
3 plt.plot(fpr, tpr)
4 plt.plot(fpr,fpr, '--',color = 'red')
5 plt.title('ROC curve')
6 plt.xlabel('FPR')
7 plt.ylabel('TPR')
8 plt.show()
9
10 ## AUROC score
11 roc_auc_score(y_val, y_val_pred)
```



Out[243]: 0.8218876305413652

```
In [162]: 1 # Check PR curve and AU PR curve value
2 pr, rc, th = precision_recall_curve(y_test,y_test_pred1)
3 plt.plot(rc,pr)
4 plt.title('Pr curve')
5 plt.xlabel('Recall')
6 plt.ylabel('Precision')
7 plt.show()
8
9 # AUC score of PR curve
10 print('AUC score of PR curve :',auc(rc, pr))
```




AUC score of PR curve : 0.6738161213579151

-----

Insights :

- After reducing dimensions to 2 feature to view data. Data as more misclassification. Classification data is difficult.
- Test F1 score is 0.67 which low and as per data this is possible.
- Recall and precision should be balanced while dealing with banking domain. in R curve we can pick where precison is high and even recall is high. By above PR curve Precis ion can 0.65 and recall can be 0.68.
- - With removal of 'installment', 'sub\_grade','grade' from data, VIF value got reduced below 5. So there are columns with very less multicollinearity between independent c olumns.





Recommendations

- As per weight data from model, emp\_title, title, int\_rate and term as high importance. So while providing loan to a person these feature should be noted as priority.
- With this distribution of data, KNN algorithm can be used, as KNN is a non-parametric model.

Questionnaire

- What percentage of customers have fully paid their Loan Amount? -> 80% of customers fully paid loan.
- Comment about the correlation between Loan Amount and Installment features. -> 0.042407, very less correlated.
- The majority of people have home ownership as \_\_\_\_\_. -> MORTGAGE as house ownership for majority of people.
- People with grades 'A' are more likely to fully pay their loan. (T/F) -> F, B grade people are more likely to pay loans.
- Name the top 2 afforded job titles. -> Teacher of about 4389 counts.
- Thinking from a bank's perspective, which metric should our primary focus be on? -> F1 score (ROC AUC, Precision, Recall, F1 Score)
- How does the gap in precision and recall affect the bank? Precision and recall should be balanced. If the gap between precision and recall is more than either False Positive or False Negative is more. In a bank, losses in revenue occur if a person who can repay a loan is misclassified as a defaulter and vice versa.
- Which features heavily affected the outcome? -> emp\_title and title as primary focus
- Will the results be affected by geographical location? (Yes/No) Yes, in metro cities expenses increase when compared to non-metro cities. So loan amount requirement is more in metro cities.

In [245]:

1	df
---	----

Out[245]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ownership	annual_inc	...	open_acc	pub_rec	revol_bal	revol_util	total_acc	initial_list_status	application_type	mort_acc
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	RENT	117000.0	...	16.0	0.0	36369.0	41.8	25.0	w	INDIVIDUAL	0.0
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	MORTGAGE	65000.0	...	17.0	0.0	20131.0	53.3	27.0	f	INDIVIDUAL	3.0
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	RENT	43057.0	...	13.0	0.0	11987.0	92.2	26.0	f	INDIVIDUAL	0.0
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years	RENT	54000.0	...	6.0	0.0	5472.0	21.5	13.0	f	INDIVIDUAL	0.0
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years	MORTGAGE	55000.0	...	13.0	0.0	24584.0	69.8	43.0	f	INDIVIDUAL	1.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
396025	10000.0	60 months	10.99	217.38	B	B4	licensed bankere	2 years	RENT	40000.0	...	6.0	0.0	1990.0	34.3	23.0	w	INDIVIDUAL	0.0
396026	21000.0	36 months	12.29	700.42	C	C1	Agent	5 years	MORTGAGE	110000.0	...	6.0	0.0	43263.0	95.7	8.0	f	INDIVIDUAL	1.0
396027	5000.0	36 months	9.99	161.32	B	B1	City Carrier	10+ years	RENT	56500.0	...	15.0	0.0	32704.0	66.9	23.0	f	INDIVIDUAL	0.0
396028	21000.0	60 months	15.31	503.02	C	C2	Gracon Services, Inc	10+ years	MORTGAGE	64000.0	...	9.0	0.0	15704.0	53.8	20.0	f	INDIVIDUAL	5.0
396029	2000.0	36 months	13.61	67.98	C	C2	Internal Revenue Service	10+ years	RENT	42996.0	...	3.0	0.0	4292.0	91.3	19.0	f	INDIVIDUAL	NaN

396030 rows × 27 columns

```
In [241]: 1 a = pd.DataFrame(X_res.columns)
2 b = pd.DataFrame(model1.coef_.T).rename(columns = {0: 'values'})
3 c = pd.concat([a,b], axis = 1)
4 c.sort_values(by = 'values', ascending = False)
```

Out[241]:

		0	values
20	emp_title	1.639037	
21	title	0.563966	
2	int_rate	0.421393	
1	term	0.272908	
12	revol_util	0.216163	
8	dti	0.162070	
9	open_acc	0.161318	
4	home_ownership	0.143101	
0	loan_amnt	0.115683	
5	annual_inc	0.087177	
6	verification_status	0.027544	
15	application_type	0.001926	
10	pub_rec	0.000000	
17	pub_rec_bankruptcies	0.000000	
13	total_acc	-0.038575	
19	earliest_cr_line_diff	-0.041096	
11	revol_bal	-0.048602	
16	mort_acc	-0.069787	
3	emp_length	-0.116052	
18	issue_d_diff	-0.139975	
14	initial_list_status	-0.143197	
7	purpose	-0.194034	

```
In [202]: 1 type(model1.coef_)
```

Out[202]: numpy.ndarray

```
In [191]: 1 X_train.shape
```

Out[191]: (382201, 22)

```
In [190]: 1 model1.coef_.shape
```

Out[190]: (1, 22)

```
In [174]: 1 df['emp_title'].value_counts()
```

Out[174]: Teacher 4389  
Manager 4250  
Registered Nurse 1856  
RN 1846  
Supervisor 1830  
...  
Postman 1  
McCarthy & Holthus, LLC 1  
jp flooring 1  
Histology Technologist 1  
Gracon Services, Inc 1  
Name: emp\_title, Length: 173105, dtype: int64

```
In [171]: 1 df.groupby(['grade'])['loan_status'].count().reset_index().sort_values(by= 'loan_status', ascending = False)
```

Out[171]:

	grade	loan_status
1	B	116018
2	C	105987
0	A	64187
3	D	63524
4	E	31488
5	F	11772
6	G	3054